




Session Processing

Greg Utas
Pentennea
greg@pentennea.com

Personal Background

- ❖ software architecture
- ❖ Nortel (20 years)
 - call server frameworks
 - GSM MSC rearchitecture
- ❖ Sonim (2 years)
 - push-to-talk for wireless networks
- ❖ Pentennea
 - carrier-grade software consulting



Session Processing Patterns



- ❖ Definitions
- ❖ Protocol
- ❖ State Machine
- ❖ Separation of Services and I/O
- ❖ Separation of Services and Protocols
- ❖ Separation of Sessions
- ❖ Message Preservation
- ❖ Type-Length-Value Message

Definitions



| Connectionless protocol | Connection-oriented protocol |
|-------------------------|----------------------------------|
| stateless | stateful |
| request-response | setup, data exchange, disconnect |
| transaction processing | session processing |

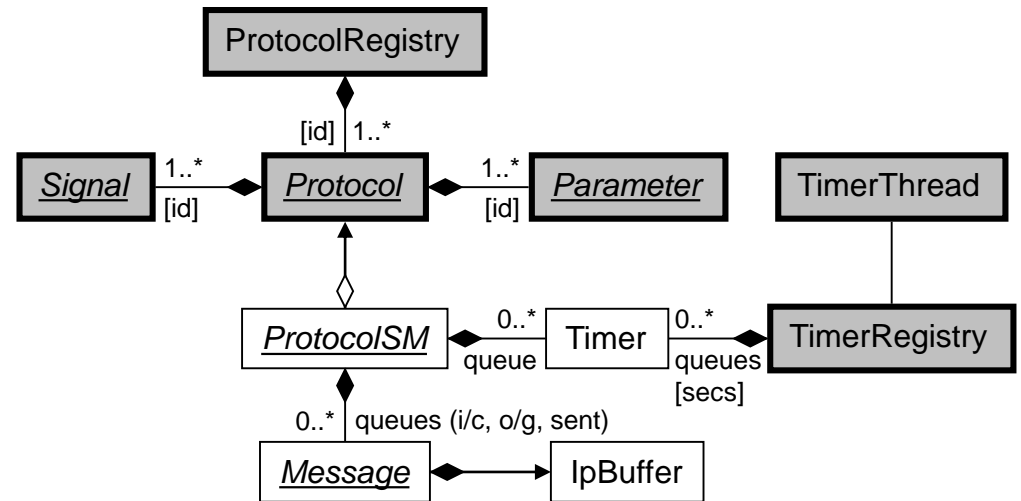
- ❖ transaction: receive message, perform work, send message(s)
 - handling a request in a connectionless protocol
 - handling a message in a connection-oriented protocol
- ❖ session: user-network (client-server) dialog
 - based on a connection-oriented protocol
- ❖ service: uses protocols to implement a user application
 - may connect users (data exchange phase = user-to-user communication)

Protocol



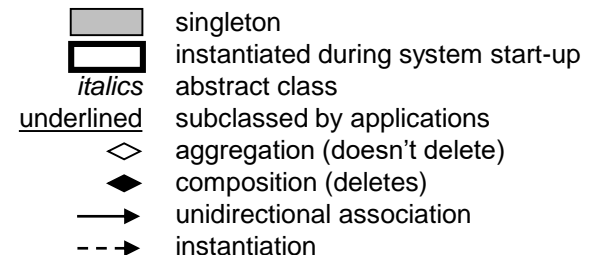
- ❖ a Protocol defines
 - the order in which Signals (message types) can be sent
 - the Parameters that are mandatory or optional for each signal
- ❖ a Message
 - provides functions for parsing and building a byte stream
 - owns an IpBuffer, which contains the byte stream and the source and destination IP addresses
- ❖ each message passes through a ProtocolSM (PSM) that
 - implements a state machine to enforce its protocol
 - is created (destroyed) during the transaction in which an initial (final) message in its protocol is sent or received
 - communicates with a conjugate PSM (if an internal protocol) or the outside world (if an external protocol)

Object Model (for Protocols)



❖ each Timer is associated with a PSM

- timeout message to PSM when timer expires
- timer placed in registry (timewheel) serviced by timer thread
- granularity = 1 second (soft real-time)

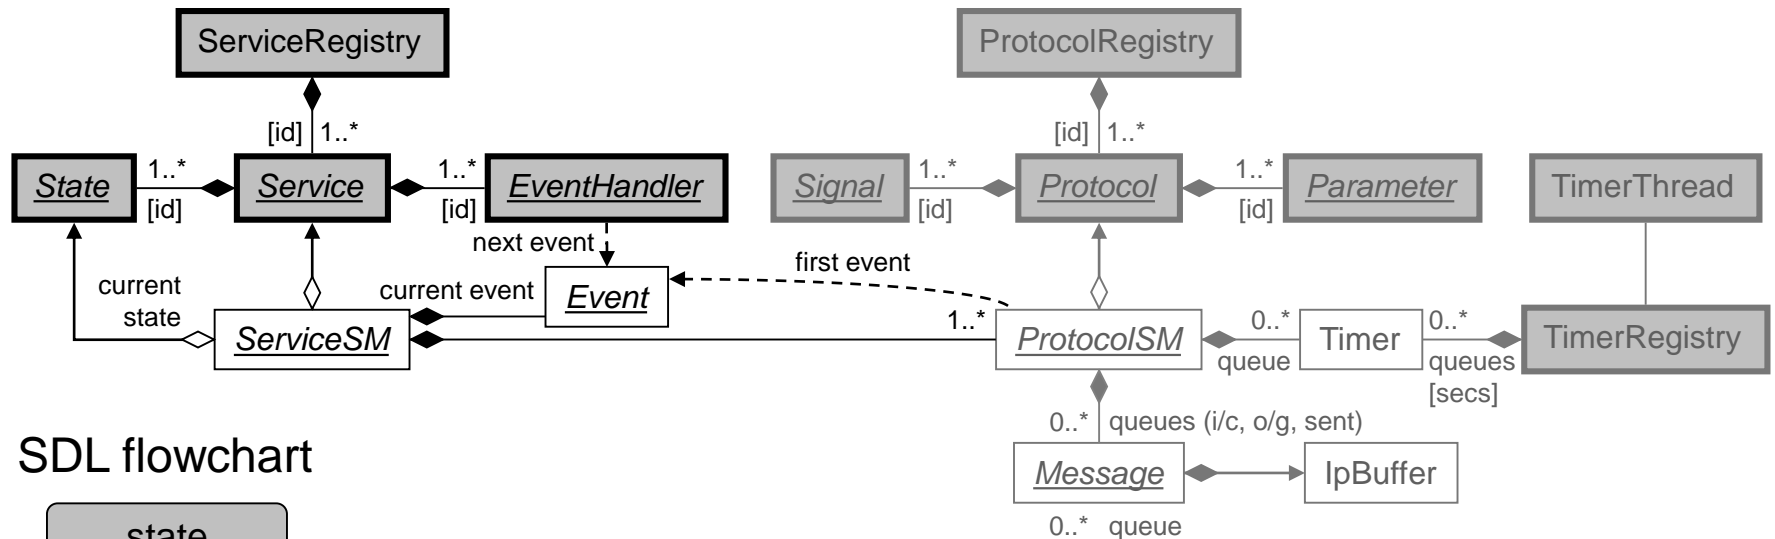


Service




- ❖ a Service defines a set of
 - Events: specify work that the service should perform
 - are objects (to allow arguments)
 - States: identify the event handler associated with an event
 - EventHandlers: react to an event in some state
 - perform work, including building outgoing message(s)
 - determine the next state and event
- ❖ a ServiceSM (SSM) is a per-user instance of a service
 - has a current state and event, and a next state and event
- ❖ invocation of event handlers done within framework
 - confines application software to event handlers and SSM functions (precludes service logic in state machine drivers)
 - simplifies patching of missing event handlers

Object Model (adding Services)



SDL flowchart

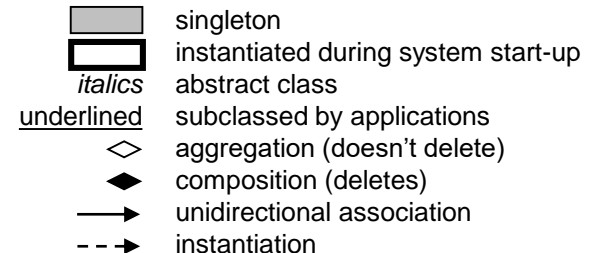


void ServiceSM::ProcessEvent(Event* currEvent, Event* &nextEvent):

```

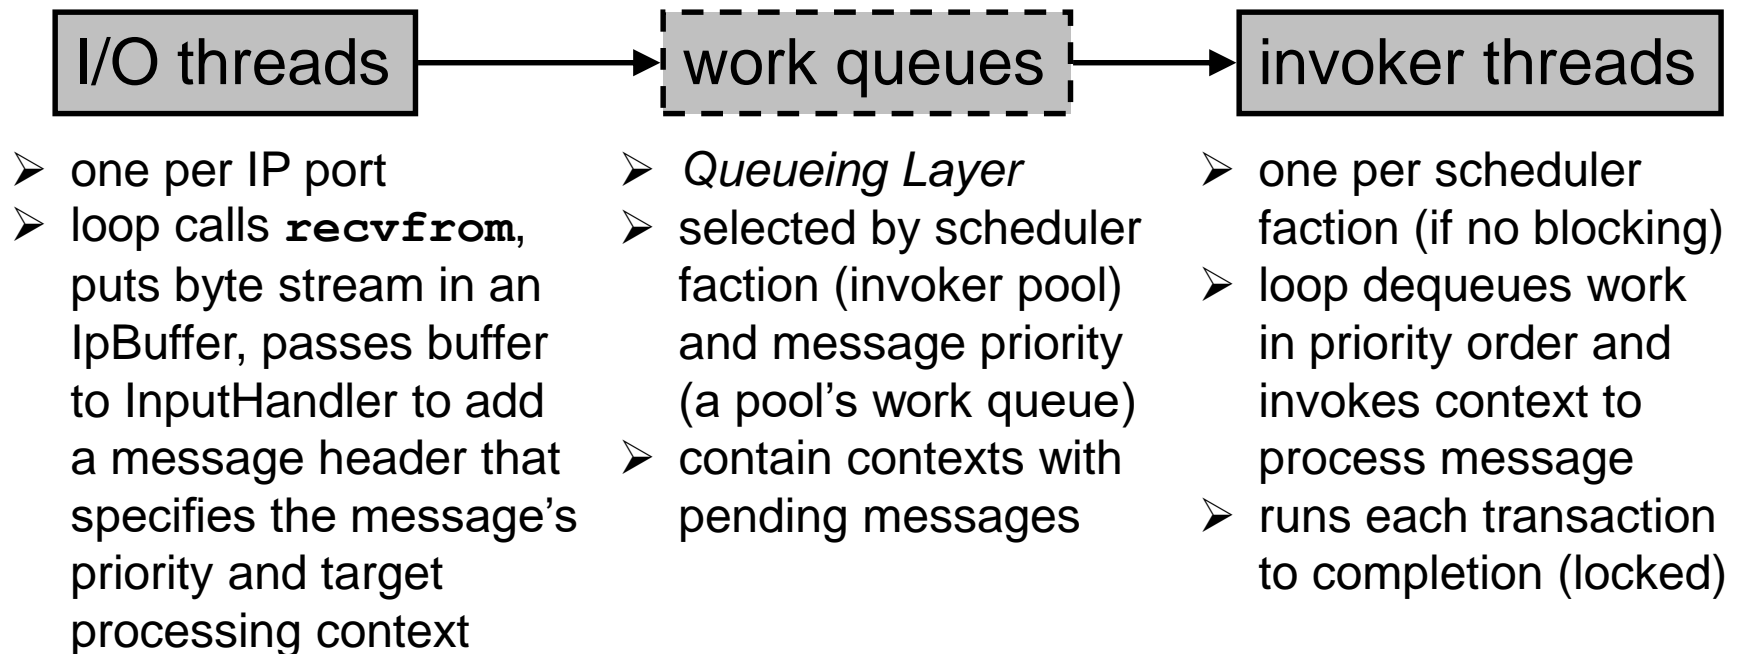
...
Event::Id eid = currEvent->Eid();
EventHandler::Id ehid = currState->Handler(eid);
EventHandler* handler = service->Handler(ehid);
handler->ProcessEvent(*this, *currEvent, nextEvent);
...
    
```

event handler



Separation of Services and I/O

- ❖ use *Half-Sync/Half-Async* pattern to
 - eliminate I/O blocking in services
 - allow different services to use the same protocol
 - prioritize incoming work (for overload controls)





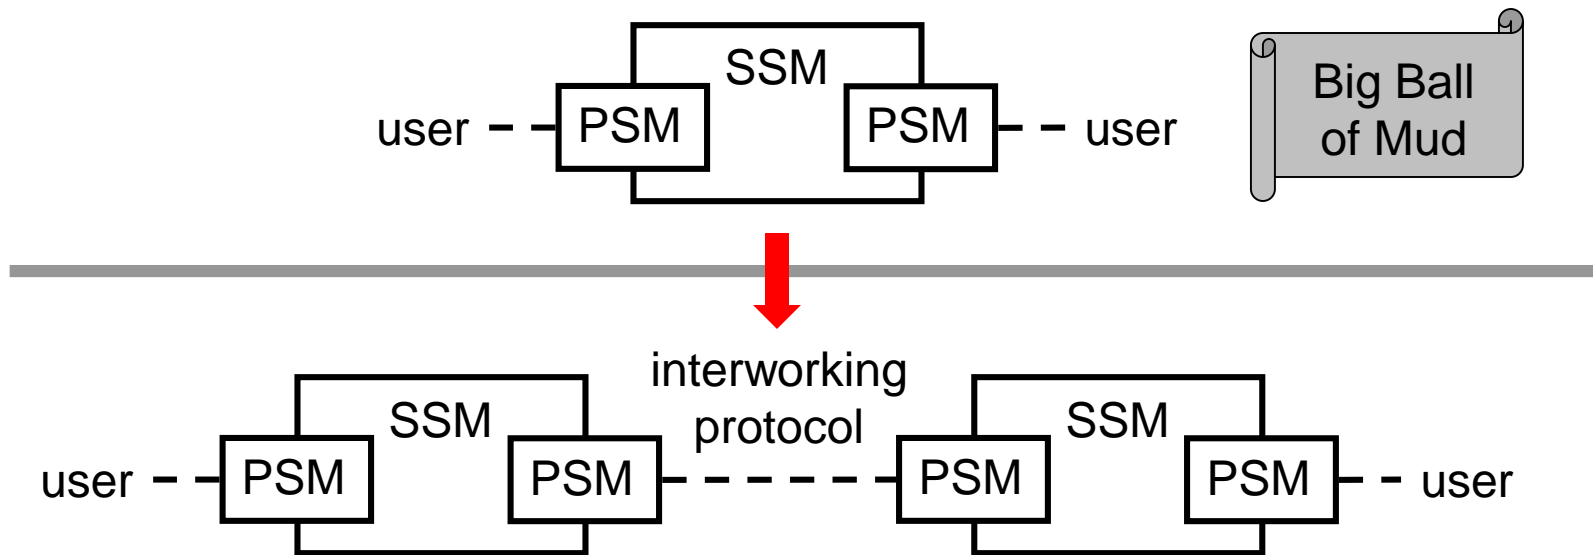
Separation of Services & Protocols



- ❖ many services use more than one protocol simultaneously
 - consequently, service states usually differ from protocol states
 - therefore have separate ServiceSM and ProtocolSM classes
- ❖ must map protocol input (signal) to service input (event)
 - implemented by PSM raising first event for SSM
 - usually, AnalyzeMessage event
 - SSM does `switch(signal)` to determine the actual event
 - rarely, ProtocolError event
 - when PSM detects a violation that requires SSM involvement
 - localizes code changes when enhancing or replacing a protocol
 - is a cleaner solution for *stimulus protocols*, in which the meaning of a signal depends on state
 - offhook signal maps to origination, answer, or reanswer event
 - mouse clicks and soft keys are also state-dependent

Separation of Sessions

- ❖ use separate sessions when connecting users
 - scalability: users can run on different processors
 - requires a database that maps users to processors
- ❖ sessions communicate using an interworking protocol
 - avoids $O(n^2)$ development cost to interwork n user-network protocols





Message Preservation

- ❖ to reduce memory usage, default behavior is that
 - an incoming message is deleted at the end of the transaction in which it is processed
 - an outgoing message is deleted after it is sent
- ❖ however, the ability to save messages is useful
 - to reference parameters in an incoming message during a subsequent transaction
 - to retransmit an outgoing message if an acknowledgment does not arrive
- ❖ implemented by **save** and **unsave** functions
 - increment and decrement a counter
 - delete message when counter drops to zero

Type-Length-Value Message



| Technique | Synopsis |
|----------------------|---|
| TLV Message | type (parameter identifier) + length + value (contents)—more efficient than text, ASN.1 or XML encodings |
| Parameter Typing | cast contents as a struct to improve reliability and readability |
| Parameter Fence | AddParm places pattern (e.g. 0xaaaaaaaa) after parameter to detect trampling when next parameter is added |
| Parameter Template | AddParm uses parameter identifier to access a template that initializes the parameter |
| Parameter Dictionary | parse a message once, when it arrives, and construct a lookup table for fast access to its parameters; table is used by FindParm |

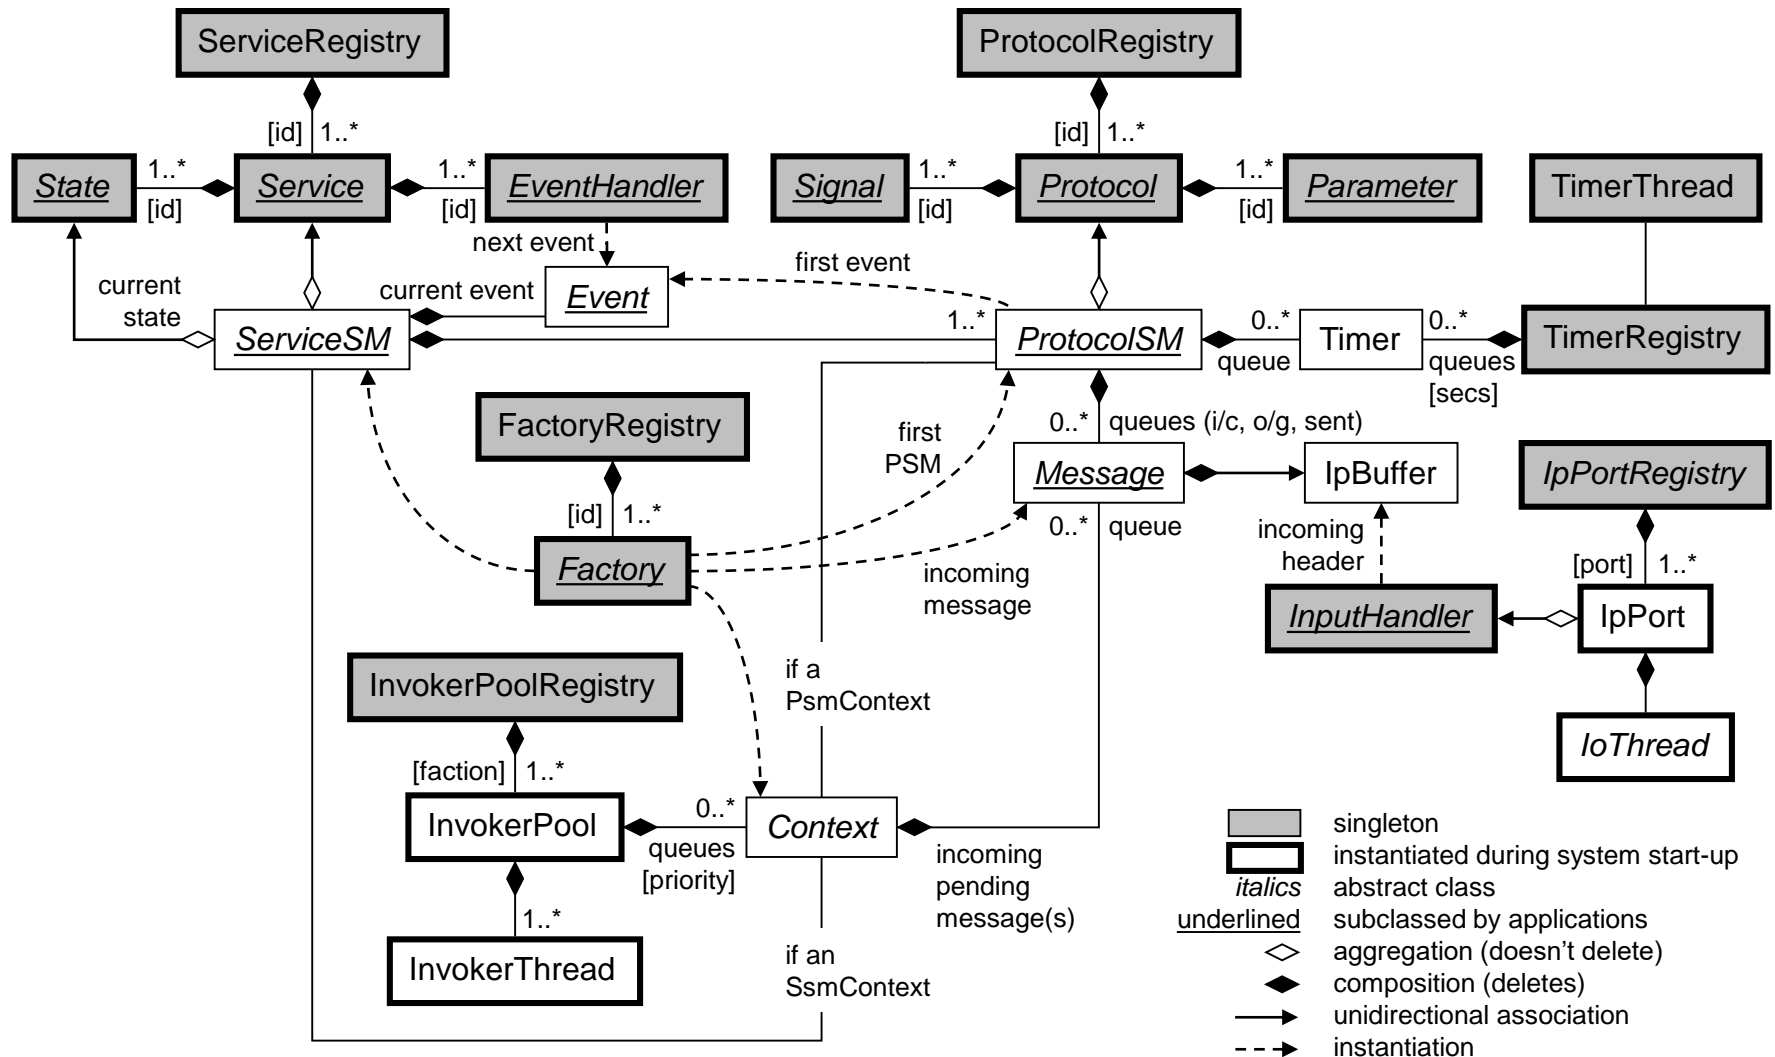
❖ implement TlvMessage as a subclass of Message

Session Processing Framework

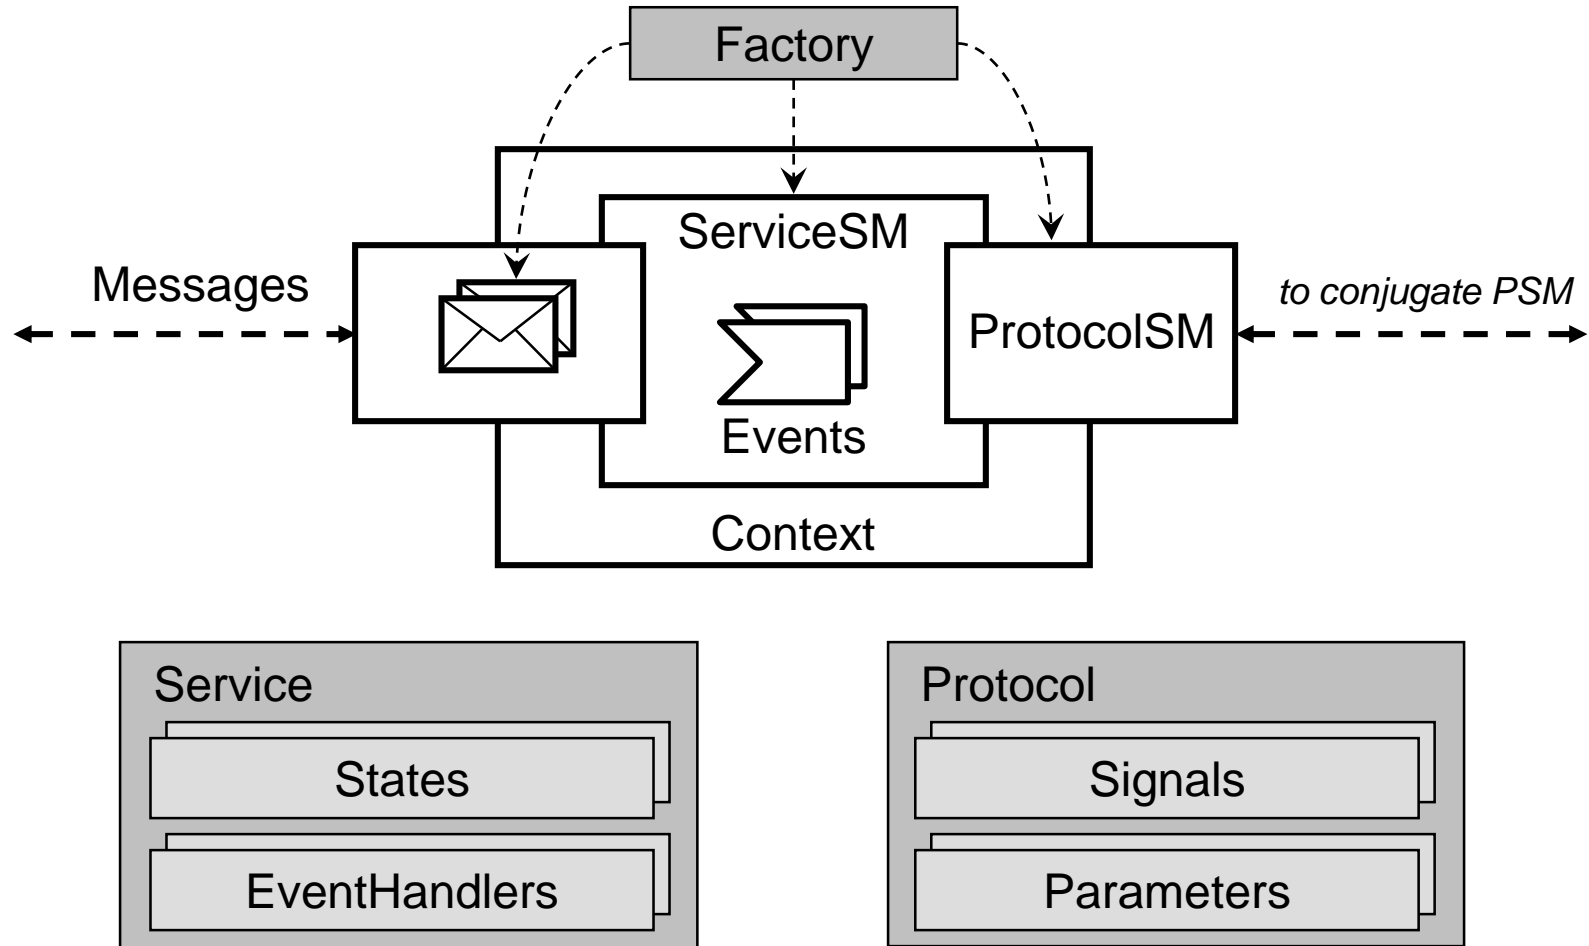


- ❖ Object Model
- ❖ Run-Time View
- ❖ Class Responsibilities
- ❖ Application Subclasses
- ❖ Application Module Structure
- ❖ Changes in Terminology

Object Model



Run-Time View



Class Responsibilities



| Class | Purpose |
|------------------|---|
| Service | provide registries for a service's states and event handlers |
| State | define the event handler to be invoked for each event that is legal in this state |
| Event | define an input to a service |
| EventHandler | handle one or more state-event pairs that can arise in a service |
| ServiceSM (SSM) | provide per-context data for a service that coordinates the behavior of PSMs |
| Protocol | provide registries for a protocol's signals and parameters |
| Signal | define a message type in a protocol |
| Parameter | provide additional information for one or more signals |
| IpBuffer | wrap a message (byte stream) and store source/destination IP addresses |
| InputHandler | add a header to an IpBuffer so that an incoming message can be routed |
| Message | manage an IpBuffer and provide functions to parse or build its contents |
| Timer | inject a timeout message to its PSM if it expires |
| ProtocolSM (PSM) | implement the state machine for the initiator or recipient role in a protocol |
| Factory | create incoming messages and the initial PSM, SSM, and context |
| Context | provide a scheduling <i>Façade</i> for a session's objects (messages, PSMs, SSM) |

Application Subclasses




- ❖ Service, ServiceSM, and Factory: one per service
 - State: one per state
 - Event: one per event
 - EventHandler: one per state-event combination
- ❖ Protocol: one per protocol
 - Signal: one per signal
 - Parameter: one per parameter
 - Message: one per protocol, or even one per signal
 - ProtocolSM: one per protocol role (client and server, or initiator and recipient)

Application Module Structure



- ❖ services use protocols, but not vice versa
 - if S_1 and S_2 communicate using P , can then build S_1 and P into one load and S_2 and P into another



evh EventHandlers
fac Factory
msg Messages
pro Protocol, Signals, Parameters
psm ProtocolSM
ssm ServiceSM
svc Service, States, Events
→ #includes relationship
 (transitive #includes not shown)

Changes in Terminology



| <i>A Pattern Language of Call Processing</i> | <i>Robust Services Core</i> |
|--|---|
| FSM | Service |
| PSM | ProtocolSM (PSM) |
| AFE | ServiceSM (SSM) |
| Agent | Factory |
| Transactor | Context |
| State, Event, EventHandler, Message | same |
| PFE, PFQ | not discussed (used to decouple supplementary services) |
| not discussed | other classes in object model |

Session Processing Details



- ❖ Incoming Message Walkthrough
- ❖ Message Header
- ❖ Context Subclasses
- ❖ Context Determination
- ❖ Message Delivery Scenarios
- ❖ Message Crossing
- ❖ Transaction Walkthrough
- ❖ Database Design Guidelines
- ❖ Extreme Session Processing

-
- The diagram illustrates the PSM architecture with the following components and relationships:
- Factory** (grey box) is associated with **InvokerPoolRegistry** (grey box) via a dashed line labeled **[id]** with multiplicity **1..***.
 - InvokerPoolRegistry** (grey box) is associated with **InvokerPool** (white box) via a solid line labeled **[faction]** with multiplicity **1..***.
 - InvokerPool** (white box) is associated with **InvokerThread** (white box) via a solid line with multiplicity **1..*** at the **InvokerThread** end.
 - InvokerPool** (white box) is associated with **Context** (white box) via a solid line labeled **queues** with multiplicity **0..*** at the **Context** end and **[priority]** at the **InvokerPool** end.
 - Context** (white box) is associated with **Message** (white box) via a solid line labeled **incoming pending message(s)**.
 - Context** (white box) is associated with **Message** (white box) via a dashed line labeled **incoming message**.
 - Context** (white box) is associated with **Message** (white box) via a dashed line labeled **if a PsmContext**.
 - Context** (white box) is associated with **Message** (white box) via a dashed line labeled **if an SsmContext**.
 - Message** (white box) is associated with **IpBuffer** (white box) via a solid line labeled **queue** with multiplicity **0..*** at the **Message** end.
 - Message** (white box) is associated with **IpBuffer** (white box) via a solid line labeled **incoming header**.
 - Message** (white box) is associated with **InputHandler** (grey box) via a solid line labeled **incoming header**.
 - InputHandler** (grey box) is associated with **IpPort** (white box) via a solid line with multiplicity **1..*** at the **IpPort** end.
 - IpPort** (white box) is associated with **IoThread** (white box) via a solid line with multiplicity **1..*** at the **IpPort** end.
 - IpPortRegistry** (grey box) is associated with **IpPort** (white box) via a solid line labeled **[port]** with multiplicity **1..***.

Message Header

- ❖ interprocessor messages transported by UDP (IP address + port)
- ❖ sender supplies message header for all internal messages
- ❖ InputHandler supplies message header when external message arrives

```

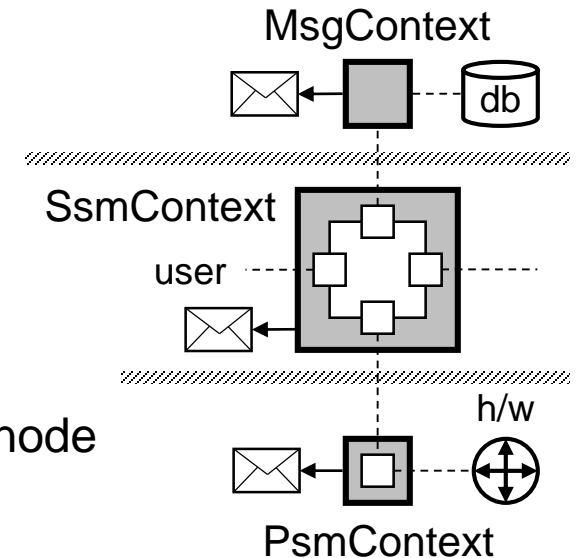
struct MessageHeader
{
    LocalAddress    txaddr;        // source address
    LocalAddress    rxaddr;        // destination address
    unsigned int    priority : 2;  // ingress, egress, progress, immediate
    unsigned int    initial : 1;   // set if first message in protocol
    unsigned int    final : 1;     // set if last message in protocol
    unsigned int    join : 1;      // set to create PSM and join existing context
    unsigned int    protocol : 16; // protocol identifier (indexes ProtocolRegistry)
    unsigned int    signal : 16;   // signal identifier (indexes Protocol)
    unsigned int    length : 16;   // length of byte stream (payload)
};

struct LocalAddress
{
    unsigned int    pid : 32;      // PSM identifier (indexes PSM object pool)
    unsigned char    seqno;        // PSM incarnation number (see "Message Crossing")
    unsigned char    fid;          // factory identifier (indexes FactoryRegistry)
};
    
```


Context Subclasses



- ❖ a Context hides application objects from an invoker thread
 - incoming message queued on Context
 - Context placed on work queue
- ❖ three subclasses of Context
 - SsmContext has an SSM and PSMs
 - for a session
 - PsmContext has a single PSM
 - to access low-level functions in a remote node
 - example: controlling hardware
 - MsgContext has a single message
 - for a connectionless protocol (stateless request-response)
 - example: querying a database
- ❖ parallel Factory subclasses (SsmFactory, PsmFactory, MsgFactory) create corresponding Context subclass





Context Determination


- ❖ incoming messages go through PSMs
 - but external protocols don't supply a destination PSM address!
- ❖ route external message to the correct PSM by using a protocol-specific parameter as a key
 - key could be source IP address and port, or user identifier
- ❖ external protocol must therefore implement a database that maps a key to a PSM
 - InputHandler looks up the PSM address in the database
 - if no PSM exists, InputHandler sets initial=true, and the Factory creates a new PSM and registers it against the appropriate key
 - PSM's destructor removes PSM from the database
- ❖ some external protocols echo a previously supplied value
 - avoid lookup overhead by getting the PSM address echoed

Message Delivery Scenarios




❖ message header determines how to deliver the message:

- initial=true, join=false, PSM=nil
 - context does not exist
 - create context, and SSM and/or PSM if required
- initial=false, join=false, PSM=non-nil
 - deliver message to existing PSM
- initial=true, join=true, PSM=nil
 - SSM exists, but PSM does not
 - create the PSM that will receive the message
 - Factory uses protocol's database to find the SSM
- initial=false, join=false, PSM=nil
 - subsequent message *before acknowledgment*
 - source PSM has not yet learned the destination PSM's address
 - destination node must find the destination PSM by using the source PSM's address as a key (search all PSMs or implement a database)



Message Crossing

- ❖ must take this type of nasty situation into account:



- ❖ calling context is deleted after it sends Release message
- ❖ problem 1: PSM (object pool block) in calling context is reallocated before Answer message arrives
 - new PSM receives a spurious Answer message
 - fix by adding an incarnation number to PSM addresses
 - discard a message that contains the wrong incarnation number
- ❖ problem 2: billing records disagree on whether call was answered
 - add “call answered” flag to Release message for correlation



Transaction Walkthrough

- ❖ InvokerThread dequeues Context from work queue
- ❖ Context dequeues Message and queues it on destination PSM
- ❖ Context invokes destination PSM to handle Message
- ❖ PSM updates its state and raises AnalyzeMessage Event
- ❖ Context invokes SSM to handle Event raised by PSM
- ❖ SSM loops, invoking EventHandlers until next event is **nullptr**
 - first EventHandler is a message analyzer that maps the incoming signal to a service-specific event (see “Separation of Services and Protocols”)
- ❖ Context invokes each PSM that has a pending outgoing message
- ❖ PSMs update their states and send their messages
- ❖ Context deletes any PSM that is in the Null state
- ❖ if SSM is in the Null state, Context deletes the SSM and itself, else Context dequeues the next Message (if any) or returns to the InvokerThread

Database Design Guidelines



- ❖ databases used during session processing
 - reside in memory
 - improves throughput
 - support non-blocking queries
 - avoiding blocking operations improves throughput; also simplifies software by eliminating critical regions
 - queries support the option to return “record locked” rather than blocking until record becomes available
- ❖ user profile is cloned when session is initiated
 - can be fetched by request-reply message sequence
 - messaging is asynchronous, so blocking is OK in this case
 - clone eliminates subsequent blocking and allows the master to be modified in parallel with the session

Extreme Session Processing



| Technique | Usage |
|-------------------------|--|
| object pools with audit | Contexts, PSMs, Messages, IpBuffers, Timers, SSMs, and Events each have their own pool |
| cooperative scheduling | I/O, invoker, and timer threads run to completion |
| asynchronous messaging | all Messages sent asynchronously |
| daemons | I/O, invoker, and timer threads created during start-up |
| proportional scheduling | I/O, invoker, and timer threads have factions |
| defensive coding | in applications generally; use Timer when waiting for ack |
| write-protected memory | most singletons (if data rarely changes after start-up) |
| safety net | I/O and invoker threads only lose current work item |
| obituaries | PSMs send/inject final messages (on death of context or node) |
| overload controls | work queues indexed by message priority |
| software error log | invoker threads can abort current work (context death) |
| object dump | display objects involved in current work (context death) |
| trace tools | tracing of messages and transactions |